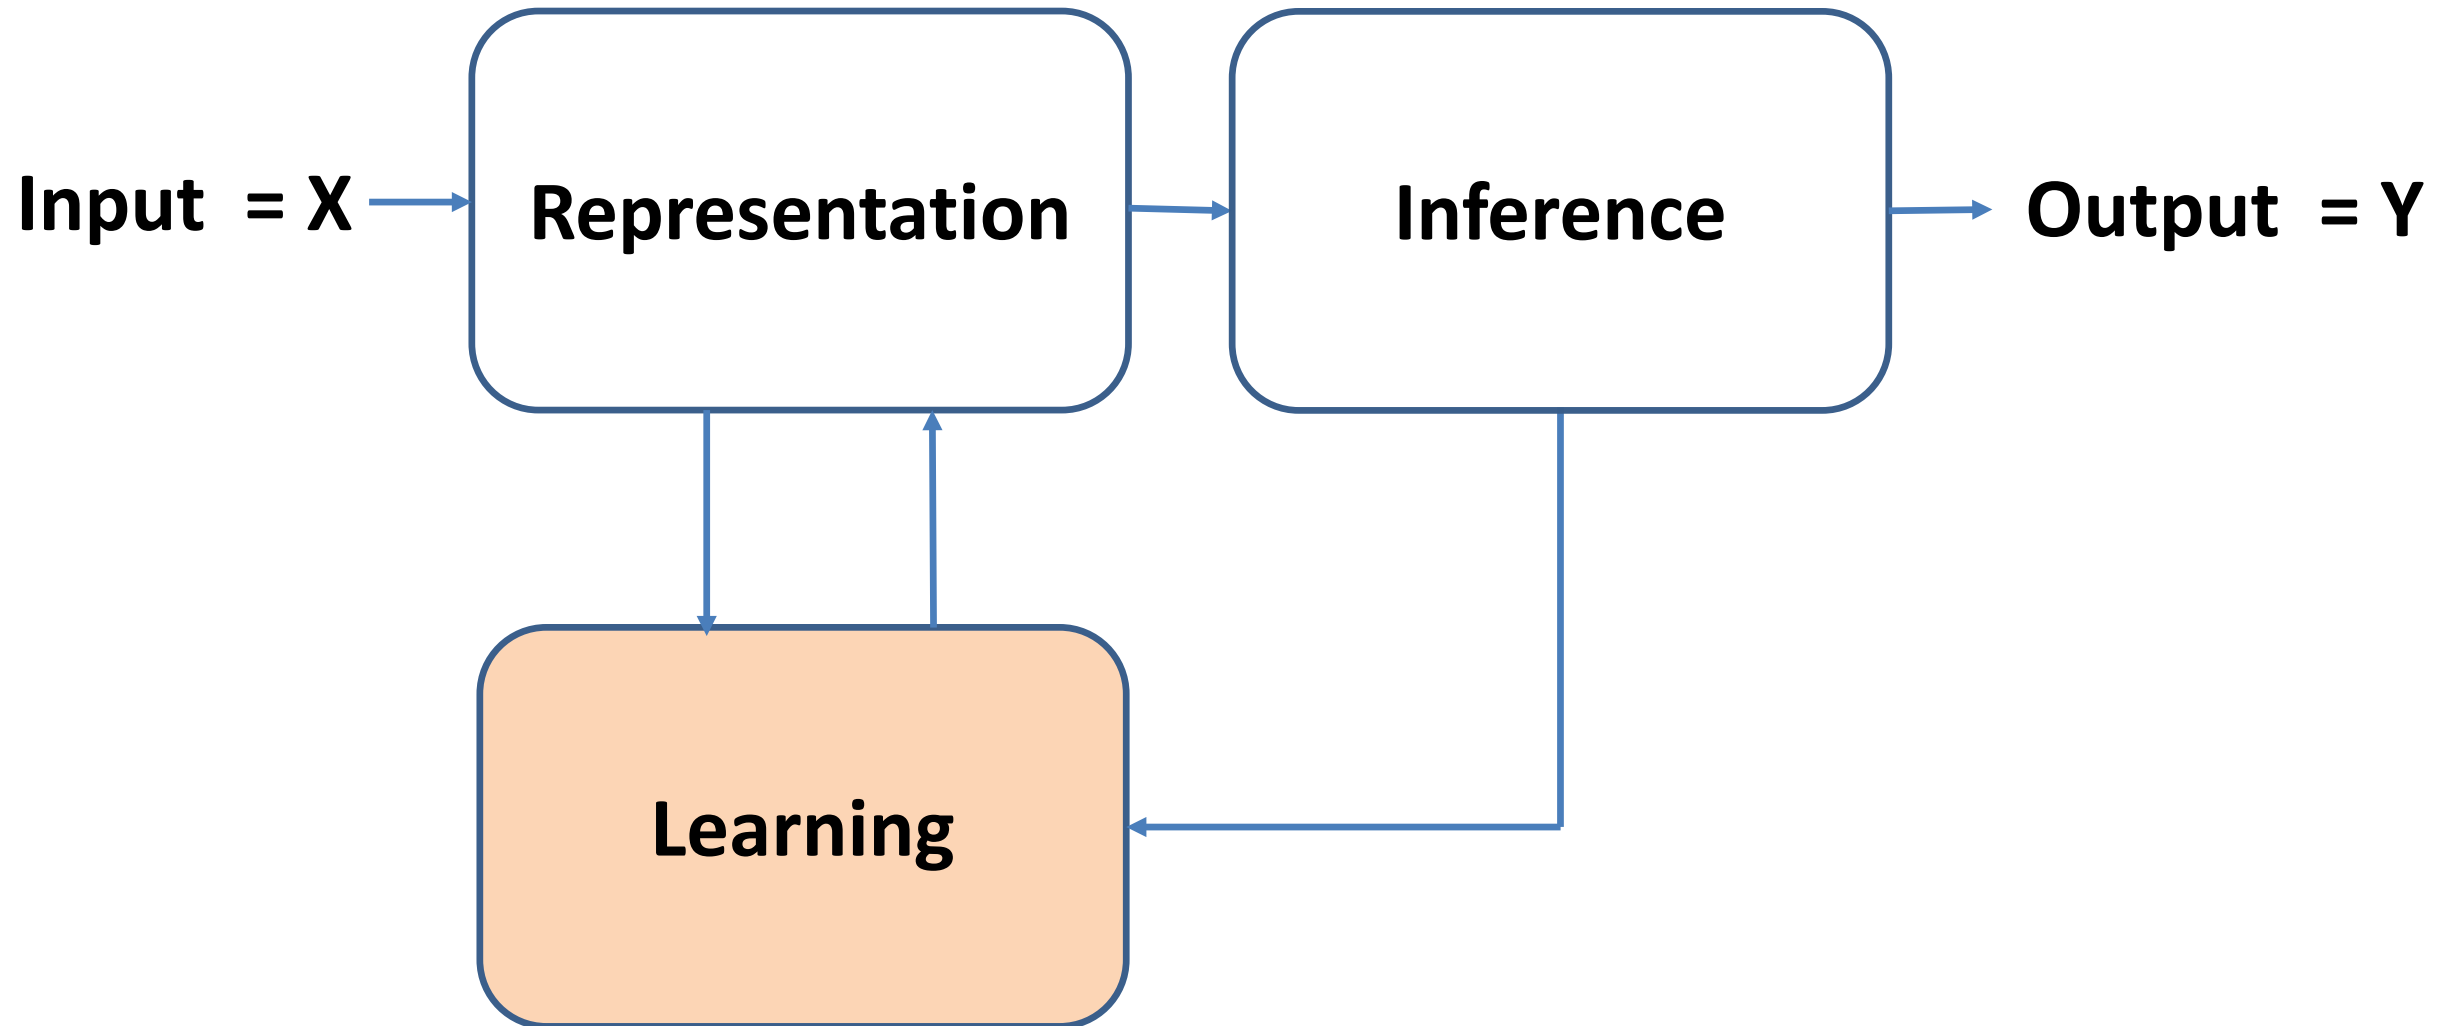


04 | Regularization



Stephen F Elston | Principle Consultant, Quantia Analytics, LLC

Introduction to Regularization for Deep Learning



Introduction to Regularization for Deep Learning

- Deep learning models have very large numbers of parameters which must be learned.
 - Even with large training datasets there may only be a few samples per parameters
- Large number of parameters leads to high chance of **over-fitting** deep learning models
 - Over-fit models do not generalize
 - Over-fit models have poor response to input noise
- To prevent over-fitting we apply **regularization methods**

Introduction to Regularization for Deep Learning

- Bias-variance trade-off
- l_2 regularization
- l_1 regularization
- Early stopping
- Dropout regularization
- Batch normalization

The Bias-Variance Trade-Off

- High capacity models fit training data well
 - Exhibit high variance
 - Do not generalize well; exhibit **brittle behavior**
 - $\text{Error}_{\text{training}} \ll \text{Error}_{\text{test}}$
- Low capacity models have high bias
 - Generalize well
 - Do not fit data well
- Regularization adds bias
 - Strong regularization adds significant bias
 - Weak regularization leads to high variance

The Bias-Variance Trade-Off

- How can we understand the bias-variance trade-off?
- We start with the error:

$$\Delta y = E[Y - \hat{f}(X)]$$

Where:

Y = the label vector.

X = the feature matrix.

$\hat{f}(x)$ = the trained model.

The Bias-Variance Trade-Off


- We can expand the error term

$$\Delta x = \left(E[\hat{f}(X)] - \hat{f}(X)\right)^2 + E\left[(\hat{f}(X) - E[\hat{f}(X)])^2\right] + \sigma^2$$

$$\Delta x = \textit{Bias}^2 + \textit{Variance} + \textit{Irreducible Error}$$

- Increasing bias decreases variance
- Notice that even if the bias and variance are 0 there is still irreducible error

The Bias-Variance Trade-Off



I2 Regularization

- Over-fit models tend to have parameters (weights) with extreme values
- One way to regularize models is to limit the values of the parameters

L2 Regularization

- One way to limit the size of the model parameters is to constrain the **L2** or **Euclidian norm**:

$$\|W\|^2 = (w_1^2 + w_2^2 + \dots + w_n^2)^{\frac{1}{2}} = \left(\sum_{i=1}^n w_i^2 \right)^{\frac{1}{2}}$$

- The regularized loss function is then:

$$J(W) = J_{MLE}(W) + \lambda \|W\|^2$$


- Where λ is the regularization hyperparameter
 - Large λ increases bias but reduces variance
 - Small λ decreases bias and increases variance

l2 Regularization

- l2 regularization goes by many names
- Is called Euclidian norm regularization
- First published by Andrey Tikhonov regularization, in late 1940s
 - Only published in English in 1977
 - Is known as **Tikhonov regularization**
- In the statistics literature the method is often called **ridge regression**
- In the engineering literature is referred to as **pre-whitening**

L2 Regularization

How can you gain some intuition about L2 regularization?



L2 regularization is considered a **soft constraint**

L1 Regularization

- Regularization can be performed with other norms
- The **L1 (min-max) norm** is another common choice
- Conceptually, L1 norm limits the sum of the absolute values of the weights:

$$\|W\|^1 = (|w_1| + |w_2| + \dots + |w_n|) = \left(\sum_{i=1}^n |w_i| \right)^1$$

- The L1 norm is also known as the **Manhattan distance** or **taxi cab distance**, since it is the distance traveled on a grid between two points.

L1 Regularization


- Given the L1 norm of the weights, the loss function becomes:

$$J(W) = J_{MLE}(W) + \alpha \|W\|_1$$

- Where α is the regularization hyperparameter
 - Large α increases bias but reduces variance
 - Small α decreases bias and increases variance
- The L1 constraint drives some weights to exactly 0
 - This behavior leads to the term **lasso regularization**

l1 Regularization

A diagram helps develop some intuition on l1 regularization:



L1 regularization is a **hard constraint** on the weights

Early Stopping


- **Early stopping** is an old and simple idea
- Stop updating the model weights before the model becomes overfit
- Early stopping is analogous to l2 regularization
- We can formulate the regularized loss function as:

$$\operatorname{argmin}_W J(W) = J(W)_{MLE} + \alpha \|W\|^2$$

- Where α is the regularization hyperparameter

Early Stopping

Early stopping has a simple geometric interpretation



Dropout regularization


- Overfit deep network models tend to suffer from a problem of co-adaptation
 - With limited training data weight tensors become adapted to the training data
 - Such a model is unlikely to generalize
- We need a way to break the co-adaptation of the weight tensor

Dropout regularization

- **Dropout regularization** is a conceptually simple method unique to deep learning
 - At each step of the gradient decent some **fraction, p , of the weights are dropped-out** of each layer
 - The result is a series of models trained for each dropout sample
 - The final model is a **geometric mean** of the individual models
- Weight values are clipped in a small range as a further regularization


Dropout regularization

- Let's look at a simple example of a network with one hidden layer:



Dropout regularization

- For $p = 0.5$ here are six of the possible samples:



Batch Normalization

- In deep neural networks there is a high chance that units in a hidden layer have a **large range of output values**
 - Causes shifts in the covariance of the output values
 - Leads to difficulty computing the gradient
 - Slows convergence
- A solution is to normalize the output of the hidden layers in the network as a batch



Microsoft

©2019 Microsoft Corporation. All rights reserved. Microsoft, Windows, Office, Azure, System Center, Dynamics and other product names are or may be registered trademarks and/or trademarks in the U.S. and/or other countries. The information herein is for informational purposes only and represents the current view of Microsoft Corporation as of the date of this presentation. Because Microsoft must respond to changing market conditions, it should not be interpreted to be a commitment on the part of Microsoft, and Microsoft cannot guarantee the accuracy of any information provided after the date of this presentation. MICROSOFT MAKES NO WARRANTIES, EXPRESS, IMPLIED OR STATUTORY, AS TO THE INFORMATION IN THIS PRESENTATION.